

Happiness Score 3

October 2, 2019

```
In [36]: import numpy as np
import pandas as pd
from sklearn.metrics import mean_squared_error
data = pd.read_csv('2015.csv')
data2 = pd.read_csv('2016.csv')
data3 = pd.read_csv('2017.csv')
```

```
In [37]: # data.append(data2, ignore_index = True)
# data.append(data3, ignore_index = True)
data.drop(['Happiness Rank', 'Country', 'Region', 'Standard Error'], axis = 1, inplace = True)
data.iloc[0]
```

```
Out[37]: Happiness Score          7.58700
Economy (GDP per Capita)         1.39651
Family                           1.34951
Health (Life Expectancy)         0.94143
Freedom                          0.66557
Trust (Government Corruption)     0.41978
Generosity                       0.29678
Dystopia Residual                 2.51738
Name: 0, dtype: float64
```

```
In [38]: data2.drop(['Happiness Rank', 'Country', 'Region', 'Lower Confidence Interval', 'Upper Confidence Interval'], axis = 1, inplace = True)
data2.iloc[0]
```

```
Out[38]: Happiness Score          7.52600
Economy (GDP per Capita)         1.44178
Family                           1.16374
Health (Life Expectancy)         0.79504
Freedom                          0.57941
Trust (Government Corruption)     0.44453
Generosity                       0.36171
Dystopia Residual                 2.73939
Name: 0, dtype: float64
```

```
In [39]: data3.drop(['Happiness.Rank', 'Country', 'Whisker.high', 'Whisker.low'], axis = 1, inplace = True)
data3.iloc[0]
```

```
Out [39]: Happiness.Score          7.537000
          Economy..GDP.per.Capita.  1.616463
          Family                   1.533524
          Health..Life.Expectancy.  0.796667
          Freedom                  0.635423
          Generosity               0.362012
          Trust..Government.Corruption. 0.315964
          Dystopia.Residual         2.277027
          Name: 0, dtype: float64
```

```
In [40]: len(data3)
```

```
Out [40]: 155
```

```
In [41]: data.columns = ['Happiness Score', 'Economy', 'Family', 'Health', 'Freedom', 'Trust', 'Generosity', 'Dystopia']
          data2.columns = ['Happiness Score', 'Economy', 'Family', 'Health', 'Freedom', 'Trust', 'Generosity', 'Dystopia']
          data3.columns = ['Happiness Score', 'Economy', 'Family', 'Health', 'Freedom', 'Trust', 'Generosity', 'Dystopia']

          data = data.append(data2)
          print(len(data))
          data = data.append(data3)
          print(len(data))
```

```
315
```

```
470
```

```
In [42]: data.head()
```

```
Out [42]:
```

	Happiness Score	Economy	Family	Health	Freedom	Trust	Generosity	Dystopia
0	7.587	1.39651	1.34951	0.94143	0.66557	0.41978	0.29678	2.51738
1	7.561	1.30232	1.40223	0.94784	0.62877	0.14145	0.43630	2.70201
2	7.527	1.32548	1.36058	0.87464	0.64938	0.48357	0.34139	2.49204
3	7.522	1.45900	1.33095	0.88521	0.66973	0.36503	0.34699	2.46531
4	7.427	1.32629	1.32261	0.90563	0.63297	0.32957	0.45811	2.45176

```
In [43]: data_13 = data.iloc[:,0:4]
          data_13.head()
          data_22 = data.iloc[:,[0,1,2]]
          data_22.head()
          data_32 = data.iloc[:,[0,6,7]]
          data_32.head()
          data_47 = data.iloc[:,0:7]
          data_47.head()
```

```
Out [43]:
```

	Happiness Score	Economy	Family	Health	Freedom	Trust	Generosity
0	7.587	1.39651	1.34951	0.94143	0.66557	0.41978	0.29678
1	7.561	1.30232	1.40223	0.94784	0.62877	0.14145	0.43630
2	7.527	1.32548	1.36058	0.87464	0.64938	0.48357	0.34139
3	7.522	1.45900	1.33095	0.88521	0.66973	0.36503	0.34699
4	7.427	1.32629	1.32261	0.90563	0.63297	0.32957	0.45811

```
In [44]: features_13 = data_13.drop(['Happiness Score'], axis = 1)
features_22 = data_22.drop(['Happiness Score'], axis = 1)
features_32 = data_32.drop(['Happiness Score'], axis = 1)
features_47 = data_47.drop(['Happiness Score'], axis = 1)
target = data['Happiness Score']
```

```
In [45]: from sklearn.linear_model import LinearRegression
from sklearn.model_selection import cross_val_score
reg_13 = LinearRegression().fit(features_13, target)

y_pred_13 = reg_13.predict(features_13)
mean_squared_error(target, y_pred_13)
```

```
Out [45]: 0.38313241487431293
```

```
In [46]: reg_22 = LinearRegression().fit(features_22, target)

y_pred_22 = reg_22.predict(features_22)
mean_squared_error(target, y_pred_22)
```

```
Out [46]: 0.4328876930949258
```

```
In [47]: reg_32 = LinearRegression().fit(features_32, target)

y_pred_32 = reg_32.predict(features_32)
mean_squared_error(target, y_pred_32)
```

```
Out [47]: 0.9297629730704532
```

```
In [48]: reg_47 = LinearRegression().fit(features_47, target)

y_pred_47 = reg_47.predict(features_47)
mean_squared_error(target, y_pred_47)
```

```
Out [48]: 0.30197336216035153
```

```
In [49]: # MLP model
from sklearn.neural_network import MLPRegressor
mlp = MLPRegressor(hidden_layer_sizes=(100),max_iter=200000)
```

```
In [50]: #Multi-layer Perceptron is sensitive to feature scaling, so it is highly recommended
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
```

```

scaler.fit(features_13)
# Now apply the transformations to the data:
mlp_features_13 = scaler.transform(features_13)
mlp.fit(mlp_features_13, target)

Out [50]: MLPRegressor(activation='relu', alpha=0.0001, batch_size='auto', beta_1=0.9,
                        beta_2=0.999, early_stopping=False, epsilon=1e-08,
                        hidden_layer_sizes=100, learning_rate='constant',
                        learning_rate_init=0.001, max_iter=200000, momentum=0.9,
                        n_iter_no_change=10, nesterovs_momentum=True, power_t=0.5,
                        random_state=None, shuffle=True, solver='adam', tol=0.0001,
                        validation_fraction=0.1, verbose=False, warm_start=False)

In [51]: mlp_pred_13 = mlp.predict(mlp_features_13)
mean_squared_error(target, mlp_pred_13)

Out [51]: 0.3155719175515372

In [52]: scaler.fit(features_22)
# Now apply the transformations to the data:
mlp_features_22 = scaler.transform(features_22)
mlp.fit(mlp_features_22, target)
mlp_pred_22 = mlp.predict(mlp_features_22)
mean_squared_error(target, mlp_pred_22)

Out [52]: 0.4121454745983672

In [53]: scaler.fit(features_32)
# Now apply the transformations to the data:
mlp_features_32 = scaler.transform(features_32)
mlp.fit(mlp_features_32, target)
mlp_pred_32 = mlp.predict(mlp_features_32)
mean_squared_error(target, mlp_pred_32)

Out [53]: 0.8125211948250888

In [54]: scaler.fit(features_47)
# Now apply the transformations to the data:
mlp_features_47 = scaler.transform(features_47)
mlp.fit(mlp_features_47, target)
mlp_pred_47 = mlp.predict(mlp_features_47)
mean_squared_error(target, mlp_pred_47)

Out [54]: 0.16937233634659382

In [55]: #Plotting training and predicted target values for 2 features selected using Linear R
import matplotlib as mpl
from mpl_toolkits.mplot3d import Axes3D
import numpy as np

```

```

import matplotlib.pyplot as plt
plt.rcParams.update({'font.size': 20})

mpl.rcParams['legend.fontsize'] = 10


fig = plt.figure(figsize=(20,10))
ax = fig.gca(projection='3d')

z = target
z1 = y_pred_22
x = data.iloc[:,1]
y = data.iloc[:,2]

ax.plot(x, y, z, 'r', label='Original target values')
ax.plot(x, y, z1, 'b', label='Predicted target values')
ax.legend()
plt.xlabel('Economy')
plt.ylabel('Family')
plt.title('Linear Regression plot')

plt.show()

```



```

In [56]: #Plotting training and predicted target values for 2 features selected using MLP mode
import matplotlib as mpl
from mpl_toolkits.mplot3d import Axes3D
import numpy as np
import matplotlib.pyplot as plt
plt.rcParams.update({'font.size': 20})

```

```

mpl.rcParams['legend.fontsize'] = 10

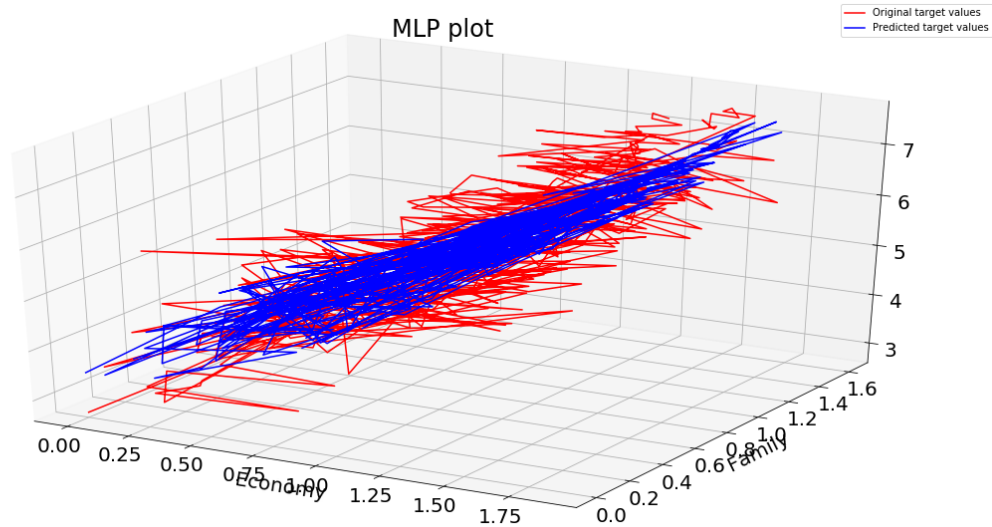
fig = plt.figure(figsize=(20,10))
ax = fig.gca(projection='3d')

z = target
z1 = mlp_pred_22
x = data.iloc[:,1]
y = data.iloc[:,2]

ax.plot(x, y, z, 'r', label='Original target values')
ax.plot(x, y, z1, 'b', label='Predicted target values')
ax.legend()
plt.xlabel('Economy')
plt.ylabel('Family')
plt.title('MLP plot')

plt.show()

```



0.1 Distribution plot between train data and predicted data


```

In [58]: import seaborn as sns
plt.figure(figsize=(10, 12))
ax1 = sns.distplot(z, hist=False, color="r", label="Actual Value(Train)")
sns.distplot(z1, hist=False, color="b", label="Fitted Values(Train)" , ax=ax1)
plt.title('Actual vs Fitted Values for Happiness score using MLP model')
plt.xlabel('Happiness score')
plt.ylabel('frequency')

```


```
plt.show()
plt.close()
```

Actual vs Fitted Values for Happiness score using MLP model



```
In [59]: import seaborn as sns
plt.figure(figsize=(10, 12))
ax1 = sns.distplot(target, hist=False, color="r", label="Actual Value(Train)")
sns.distplot(y_pred_22, hist=False, color="b", label="Fitted Values(Train)", ax=ax1)
plt.title('Actual vs Fitted Values for Happiness score using Linear regression model')
plt.xlabel('Happiness score')
plt.ylabel('frequency')
plt.show()
plt.close()
```

Actual vs Fitted Values for Happiness score using Linear regression model



```
In [63]: target = data.iloc[:,0]
         features = data.iloc[:,1:]
         print(target.head())
         features.head()
```

```
0    7.587
1    7.561
2    7.527
3    7.522
4    7.427
```

Name: Happiness Score, dtype: float64

```
Out [63]:
```

	Economy	Family	Health	Freedom	Trust	Generosity	Dystopia
0	1.39651	1.34951	0.94143	0.66557	0.41978	0.29678	2.51738
1	1.30232	1.40223	0.94784	0.62877	0.14145	0.43630	2.70201
2	1.32548	1.36058	0.87464	0.64938	0.48357	0.34139	2.49204
3	1.45900	1.33095	0.88521	0.66973	0.36503	0.34699	2.46531
4	1.32629	1.32261	0.90563	0.63297	0.32957	0.45811	2.45176


```
In [64]: reg = LinearRegression().fit(features, target)
        y_pred = reg.predict(features)
        mean_squared_error(target, y_pred)
```

```
Out[64]: 8.1869445091806e-08
```


```
In [66]: mlp = MLPRegressor(hidden_layer_sizes=(100),max_iter=200000)
        scaler = StandardScaler()
        scaler.fit(features)
        # Now apply the transformations to the data:
        mlp_features = scaler.transform(features)
        mlp.fit(mlp_features, target)
        mlp_pred = mlp.predict(mlp_features)
        mean_squared_error(target, mlp_pred)
```

```
Out[66]: 0.017458290440188224
```

```
In [65]: plt.figure(figsize=(10, 12))
        ax1 = sns.distplot(target, hist=False, color="r", label="Actual Value(Train)")
        sns.distplot(y_pred, hist=False, color="b", label="Fitted Values(Train)" , ax=ax1)
        plt.title('Actual vs Fitted Values for Happiness score using Linear regression model')
        plt.xlabel('Happiness score')
        plt.ylabel('frequency')
        plt.show()
        plt.close()
```

```
C:\Users\Welcome\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a
        return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```


Actual vs Fitted Values for Happiness score using Linear regression model(all Features)



```
In [67]: plt.figure(figsize=(10, 12))
         ax1 = sns.distplot(target, hist=False, color="r", label="Actual Value(Train)")
         sns.distplot(mlp_pred, hist=False, color="b", label="Fitted Values(Train)" , ax=ax1)
         plt.title('Actual vs Fitted Values for Happiness score using MLP model(all Features)')
         plt.xlabel('Happiness score')
         plt.ylabel('frequency')
         plt.show()
         plt.close()
```

C:\Users\Welcome\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval

Actual vs Fitted Values for Happiness score using MLP model(all Features)



In []: